

✖
 Importing necessary libraries

```

import pandas as pd # For data manipulation and analysis
import numpy as np # For numerical operations
import matplotlib.pyplot as plt # For data visualization
import seaborn as sns # For advanced data visualization
from sklearn.model_selection import train_test_split # For splitting data into training and testing sets
from sklearn.preprocessing import StandardScaler # For feature scaling
from sklearn.linear_model import LinearRegression # Linear Regression algorithm
from sklearn.ensemble import RandomForestRegressor # Random Forest algorithm
from xgboost import XGBRegressor # XGBoost algorithm
from sklearn.svm import SVR # Support Vector Regression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score # Evaluation metrics
from sklearn.impute import SimpleImputer # For handling missing values
    
```

✖
 Load the dataset

Assuming the dataset is downloaded and saved as 'Life Expectancy Data.csv'


```

df = pd.read_excel('led.xlsx') # Load the dataset
    
```

✖
 Display basic information about the dataset

```

print("Dataset Information:")
print(df.info()) # Shows columns, non-null counts, and data types
print("\nFirst 5 rows of the dataset:")
print(df.head()) # Displays first 5 rows to understand the data structure
    
```



```

5  infantdeaths      2938 non-null  int64
6  Alcohol          2744 non-null  float64
7  percentageexpenditure  2938 non-null  float64
8  HepatitisB       2385 non-null  float64
9  Measles          2938 non-null  int64
10 BMI              2904 non-null  float64
11 under-fivedeaths  2938 non-null  int64
12 Polio            2919 non-null  float64
13 Totalexpenditure  2712 non-null  float64
14 Diphtheria       2919 non-null  float64
15 HIV/AIDS         2938 non-null  float64
16 GDP              2490 non-null  float64
17 Population       2286 non-null  float64
18 thinness1-19years  2904 non-null  float64
19 thinness5-9years  2904 non-null  float64
20 Incomecompositionofresources  2771 non-null  float64
21 Schooling        2775 non-null  float64
dtypes: float64(16), int64(4), object(2)
memory usage: 505.1+ KB
None

First 5 rows of the dataset:
   Country  Year  Status  Lifeexpectancy  AdultMortality  \
0  Afghanistan  2015  Developing         65.0           263.0
1  Afghanistan  2014  Developing         59.9           271.0
2  Afghanistan  2013  Developing         59.9           268.0
3  Afghanistan  2012  Developing         59.5           272.0
4  Afghanistan  2011  Developing         59.2           275.0

   infantdeaths  Alcohol  percentageexpenditure  HepatitisB  Measles  ...  \
0              62      0.01          71.279624          65.0    1154  ...
    
```

```
thinness1-19years  thinness5-9years  Incomecompositionofresources  \
0                17.2                17.3                0.479
1                17.5                17.5                0.476
2                17.7                17.7                0.470
3                17.9                18.0                0.463
4                18.2                18.2                0.454
```

```
Schooling
0      10.1
1      10.0
2       9.9
3       9.8
4       9.5
```

```
[5 rows x 22 columns]
```

```
df.head()
```



	Country	Year	Status	Lifeexpectancy	AdultMortality	infantdeaths	Alcohol	percentageexpenditure	HepatitisB	Measles
0	Afghanistan	2015	Developing	65.0	263.0	62	0.01	71.279624	65.0	1154
1	Afghanistan	2014	Developing	59.9	271.0	64	0.01	73.523582	62.0	492
2	Afghanistan	2013	Developing	59.9	268.0	66	0.01	73.219243	64.0	430
3	Afghanistan	2012	Developing	59.5	272.0	69	0.01	78.184215	67.0	2787
4	Afghanistan	2011	Developing	59.2	275.0	71	0.01	7.097109	68.0	3013

```
5 rows x 22 columns
```

## ▼ Data Preprocessing

### Check for missing values

```
print("\nMissing values in each column:")
print(df.isnull().sum()) # Counts null values in each column
```



```
Missing values in each column:
Country          0
Year             0
Status           0
Lifeexpectancy  10
AdultMortality   10
infantdeaths     0
Alcohol          194
percentageexpenditure  0
HepatitisB       553
Measles          0
BMI              34
under-fivedeaths  0
Polio            19
Totalexpenditure  226
Diphtheria       19
HIV/AIDS         0
GDP              448
Population       652
thinness1-19years  34
thinness5-9years   34
Incomecompositionofresources  167
Schooling        163
dtype: int64
```

## ▼ Handle missing values

We'll use median imputation for numerical columns and mode for categorical

```
numerical_cols = df.select_dtypes(include=['float64', 'int64']).columns
categorical_cols = df.select_dtypes(include=['object']).columns
```

✓ Create imputers

```
num_imputer = SimpleImputer(strategy='median') # Replaces missing values with median
cat_imputer = SimpleImputer(strategy='most_frequent') # Replaces missing values with most frequent value
```

✓ Apply imputation

```
df[numerical_cols] = num_imputer.fit_transform(df[numerical_cols])
df[categorical_cols] = cat_imputer.fit_transform(df[categorical_cols])
```

✓ Verify no missing values remain

```
print("\nMissing values after imputation:")
print(df.isnull().sum())
```

```
➡
Missing values after imputation:
Country          0
Year             0
Status           0
Lifeexpectancy  0
AdultMortality  0
infantdeaths     0
Alcohol          0
percentageexpenditure  0
HepatitisB       0
Measles          0
BMI              0
under-fivedeaths 0
Polio            0
Totalexpenditure 0
Diphtheria       0
HIV/AIDS        0
GDP              0
Population       0
thinness1-19years 0
thinness5-9years 0
Incomecompositionofresources 0
Schooling        0
dtype: int64
```

✓ Exploratory Data Analysis (EDA)

Summary statistics

```
print("\nSummary statistics:")
print(df.describe()) # Shows count, mean, std, min, max, etc.
```

```
➡
Summary statistics:
      Year  Lifeexpectancy  AdultMortality  infantdeaths  Alcohol \
count  2938.000000      2938.000000      2938.000000      2938.000000  2938.000000
mean   2007.518720        69.234717       164.725664       30.303948    4.546875
std     4.613841         9.509115       124.086215       117.926501    3.921946
min    2000.000000       36.300000        1.000000        0.000000    0.010000
25%    2004.000000       63.200000       74.000000        0.000000    1.092500
50%    2008.000000       72.100000      144.000000        3.000000    3.755000
75%    2012.000000       75.600000      227.000000       22.000000    7.390000
max    2015.000000       89.000000      723.000000     1800.000000   17.870000

      percentageexpenditure  HepatitisB      Measles      BMI \
count      2938.000000      2938.000000      2938.000000      2938.000000
mean         738.251295      83.022124      2419.592240      38.381178
std        1987.914858      22.996984     11467.272489     19.935375
min           0.000000       1.000000        0.000000       1.000000
25%           4.685343      82.000000        0.000000      19.400000
50%          64.912906      92.000000       17.000000      43.500000
75%         441.534144      96.000000      360.250000      56.100000
max       19479.911610     99.000000     212183.000000     87.300000
```


	under-fivedeaths	Polio	Totalexpenditure	Diphtheria	\
count	2938.000000	2938.000000	2938.000000	2938.000000	
mean	42.035739	82.617767	5.924098	82.393125	
std	160.445548	23.367166	2.400770	23.655562	
min	0.000000	3.000000	0.370000	2.000000	
25%	0.000000	78.000000	4.370000	78.000000	
50%	4.000000	93.000000	5.755000	93.000000	
75%	28.000000	97.000000	7.330000	97.000000	
max	2500.000000	99.000000	17.600000	99.000000	

	HIV/AIDS	GDP	Population	thinness1-19years	\
count	2938.000000	2938.000000	2.938000e+03	2938.000000	
mean	1.742103	6611.523863	1.023085e+07	4.821886	
std	5.077785	13296.603449	5.402242e+07	4.397621	
min	0.100000	1.681350	3.400000e+01	0.100000	
25%	0.100000	580.486996	4.189172e+05	1.600000	
50%	0.100000	1766.947595	1.386542e+06	3.300000	
75%	0.800000	4779.405190	4.584371e+06	7.100000	
max	50.600000	119172.741800	1.293859e+09	27.700000	

	thinness5-9years	Incomecompositionofresources	Schooling
count	2938.000000	2938.000000	2938.000000
mean	4.852144	0.630362	12.009837
std	4.485854	0.205140	3.265139
min	0.100000	0.000000	0.000000
25%	1.600000	0.504250	10.300000
50%	3.300000	0.677000	12.300000
75%	7.200000	0.772000	14.100000
max	28.600000	0.948000	20.700000


## Visualize the distribution of the target variable (Life expectancy)

```
plt.figure(figsize=(10, 6))
sns.histplot(df['Lifeexpectancy'],kde=True, bins=30)
plt.title('Distribution of Life Expectancy')
plt.xlabel('Life Expectancy (years)')
plt.ylabel('Frequency')
plt.show()
```



## Correlation matrix to understand relationships between variables


```
plt.figure(figsize=(15, 10))
corr_matrix = df.corr(numeric_only=True) # Calculates correlation between numerical columns
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix of Numerical Features')
plt.show()
```



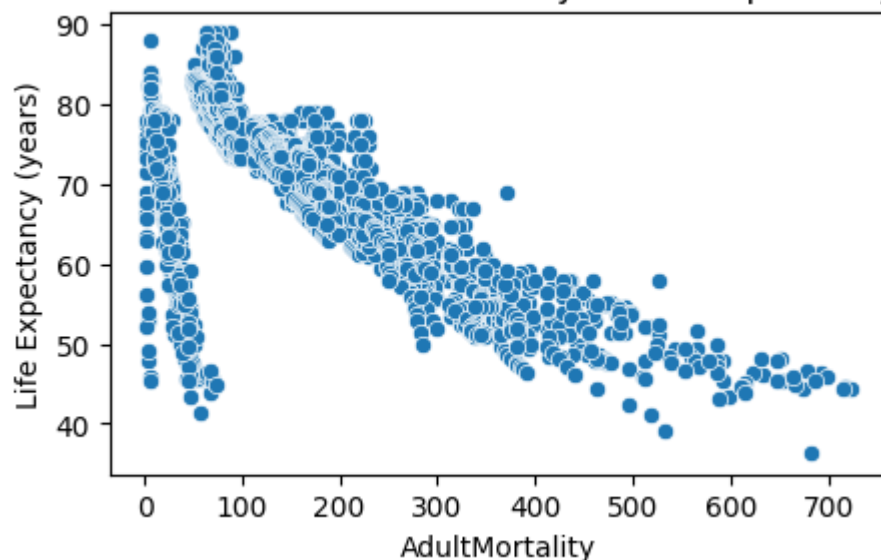
```
#plot scatter plots for each feature against Life Expectancy
for col in numerical_cols:
    if col != 'Lifeexpectancy':
        plt.figure(figsize=(5, 3))
        sns.scatterplot(x=df[col], y=df['Lifeexpectancy'])
        plt.title(f'Scatter Plot of {col} vs Life Expectancy')
        plt.xlabel(col)
        plt.ylabel('Life Expectancy (years)')
        plt.show()
```




Scatter Plot of Year vs Life Expectancy




Scatter Plot of AdultMortality vs Life Expectancy




Scatter Plot of infantdeaths vs Life Expectancy




Scatter Plot of Alcohol vs Life Expectancy




Scatter Plot of percentageexpenditure vs Life Expectancy









Scatter Plot of HepatitisB vs Life Expectancy




Scatter Plot of Measles vs Life Expectancy




Scatter Plot of BMI vs Life Expectancy




Scatter Plot of under-fivedeaths vs Life Expectancy

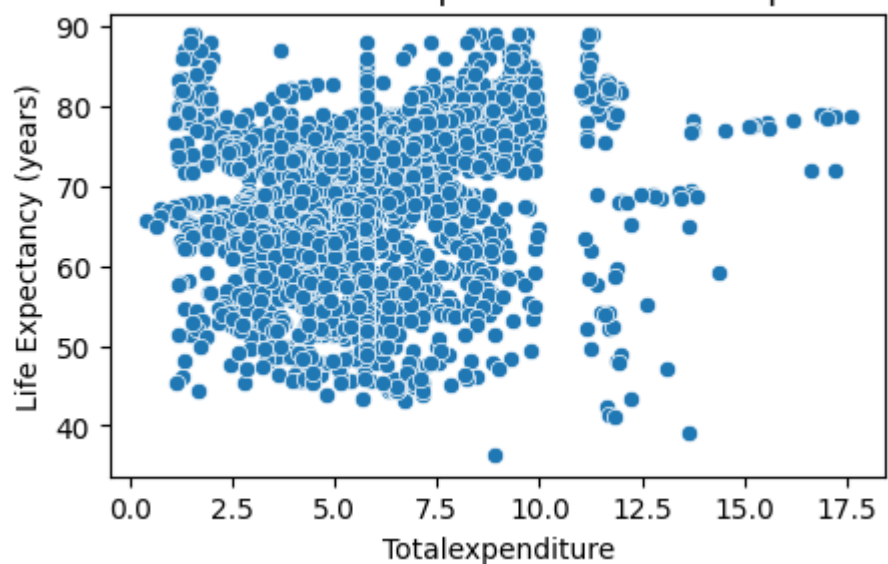


Scatter Plot of Polio vs Life Expectancy






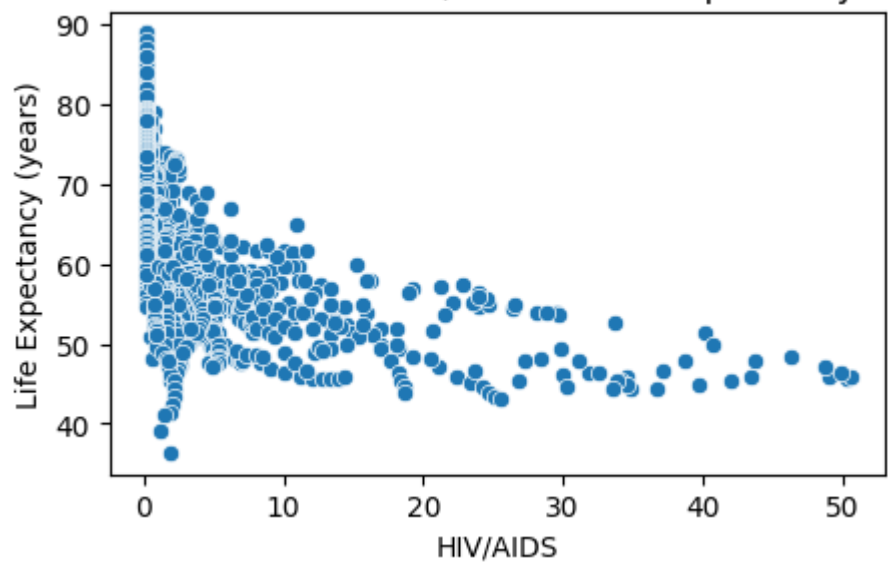
Scatter Plot of Totalexpenditure vs Life Expectancy




Scatter Plot of Diphtheria vs Life Expectancy




Scatter Plot of HIV/AIDS vs Life Expectancy




Scatter Plot of GDP vs Life Expectancy




Scatter Plot of Population vs Life Expectancy



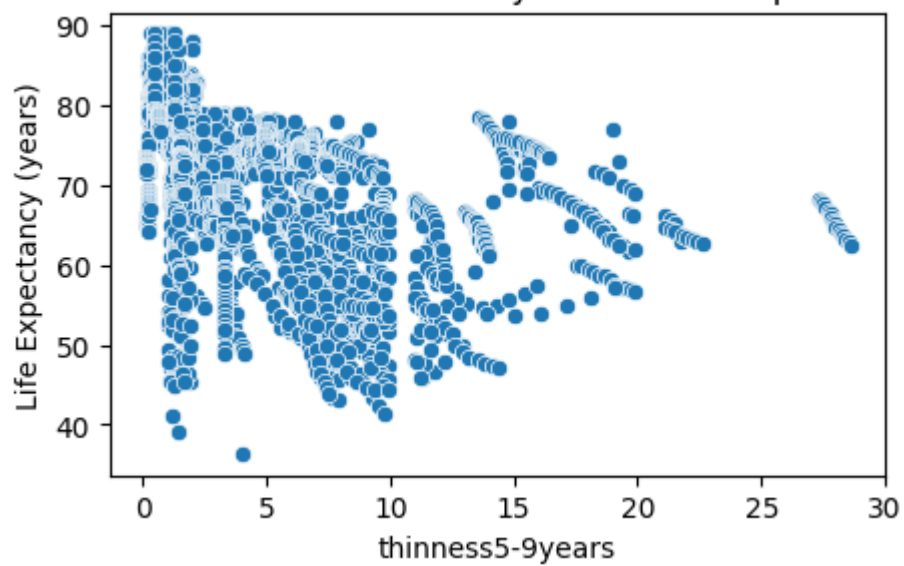





Scatter Plot of thinness1-19years vs Life Expectancy




Scatter Plot of thinness5-9years vs Life Expectancy



Scatter Plot of Incomecompositionofresources vs Life Expectancy



Scatter Plot of Schooling vs Life Expectancy



Double-click (or enter) to edit

## ✓ Feature Engineering

### Encoding categorical variables (Country and Status)

```
df = pd.get_dummies(df, columns=['Status'], drop_first=True) # Converts Status to binary (1 for Developed, 0 for Developing)
```

## ✓ For Country, since there are many unique values, we'll drop it to avoid high dimensionality

```
df.drop('Country', axis=1, inplace=True) # Removes Country column
```

## ✓ Define features (X) and target (y)

```
X = df.drop('Lifeexpectancy', axis=1) # All columns except Life expectancy
y = df['Lifeexpectancy'] # Target variable we want to predict
```

## ✓ Split the data into training and testing sets (80% train, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

## ✓ Feature Scaling

### Many algorithms perform better when features are on similar scales

```
scaler = StandardScaler() # Standardizes features by removing mean and scaling to unit variance
X_train_scaled = scaler.fit_transform(X_train) # Fit to training data and transform
X_test_scaled = scaler.transform(X_test) # Transform test data using same scaling
```

## ✓ Model Building and Evaluation

### Dictionary to store evaluation results

#### 1. Linear Regression

```
results = {}

# 1. Linear Regression
lr = LinearRegression() # Creates a linear regression model
lr.fit(X_train_scaled, y_train) # Trains the model on scaled training data
y_pred_lr = lr.predict(X_test_scaled) # Makes predictions on test data
```

## ✓ Calculate evaluation metrics

```
rmse_lr = np.sqrt(mean_squared_error(y_test, y_pred_lr)) # Root Mean Squared Error
mae_lr = mean_absolute_error(y_test, y_pred_lr) # Mean Absolute Error
r2_lr = r2_score(y_test, y_pred_lr) # R-squared score
```

```
# Store results
results['Linear Regression'] = {'RMSE': rmse_lr, 'MAE': mae_lr, 'R2': r2_lr}
```

## 2. Random Forest Regressor

```
rf = RandomForestRegressor(n_estimators=100, random_state=42) # Creates RF model with 100 trees
rf.fit(X_train_scaled, y_train) # Trains the model
y_pred_rf = rf.predict(X_test_scaled) # Makes predictions
```

Double-click (or enter) to edit

## Calculate metrics

```
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))
mae_rf = mean_absolute_error(y_test, y_pred_rf)
r2_rf = r2_score(y_test, y_pred_rf)
```

## Store results

```
results['Random Forest'] = {'RMSE': rmse_rf, 'MAE': mae_rf, 'R2': r2_rf}
```

## 3. XGBoost Regressor

```
xgb = XGBRegressor(n_estimators=100, random_state=42) # Creates XGBoost model with 100 trees
xgb.fit(X_train_scaled, y_train) # Trains the model
y_pred_xgb = xgb.predict(X_test_scaled) # Makes predictions
```

## Calculate metrics

```
rmse_xgb = np.sqrt(mean_squared_error(y_test, y_pred_xgb))
mae_xgb = mean_absolute_error(y_test, y_pred_xgb)
r2_xgb = r2_score(y_test, y_pred_xgb)
```

## Store results

```
results['XGBoost'] = {'RMSE': rmse_xgb, 'MAE': mae_xgb, 'R2': r2_xgb}
```

## 4. Support Vector Regression

```
svr = SVR(kernel='rbf') # Creates SVR model with Radial Basis Function kernel
svr.fit(X_train_scaled, y_train) # Trains the model
y_pred_svr = svr.predict(X_test_scaled) # Makes predictions
```

## Calculate metrics

```
rmse_svr = np.sqrt(mean_squared_error(y_test, y_pred_svr))
mae_svr = mean_absolute_error(y_test, y_pred_svr)
r2_svr = r2_score(y_test, y_pred_svr)
```

## Store results

```
results['Support Vector Regression'] = {'RMSE': rmse_svr, 'MAE': mae_svr, 'R2': r2_svr}
```

## ✓ Display results

```
print("\nModel Evaluation Results:")
for model, metrics in results.items():
    print(f"\n{model}:")
    print(f"RMSE: {metrics['RMSE']:.2f}") # Lower is better
    print(f"MAE: {metrics['MAE']:.2f}") # Lower is better
    print(f"R2 Score: {metrics['R2']:.2f}") # Closer to 1 is better
```



Model Evaluation Results:

Linear Regression:

RMSE: 3.91

MAE: 2.86

R2 Score: 0.82

Random Forest:

RMSE: 1.66

MAE: 1.08

R2 Score: 0.97

XGBoost:

RMSE: 1.75

MAE: 1.18

R2 Score: 0.96

Support Vector Regression:

RMSE: 3.32

MAE: 2.31

R2 Score: 0.87

## ✓ Feature Importance for the best performing model (Random Forest or XGBoost)

Let's check feature importance from Random Forest


```
feature_importance = rf.feature_importances_ # Gets importance scores from trained RF model
features = X.columns # Gets feature names
```

## ✓ Create a dataframe for visualization

```
importance_df = pd.DataFrame({'Feature': features, 'Importance': feature_importance})
importance_df = importance_df.sort_values('Importance', ascending=False) # Sorts by importance
```

## ✓ Plot feature importance

```
plt.figure(figsize=(12, 8))
sns.barplot(x='Importance', y='Feature', data=importance_df)
plt.title('Feature Importance from Random Forest')
plt.show()
```



## ✓ gdp per capita

```
#implement simple linear regression to predict life expectancy based solely on GDP per capita
df['gdp_per_capita'] = df['GDP'] / df['Population'] # Calculate GDP per capita
#MULTIPLY GDP per capita by 1000 to scale it
df['gdp_per_capita'] *= 1000000 # Scale GDP per capita for better interpretability
print("\nGDP per Capita:")
print(df['gdp_per_capita']) # Display first few rows of GDP per capita
```



```
GDP per Capita:
0      17.318314
1    1870.360746
2     19.908962
3    181.218991
4     21.331247
...
2933    35.559872
2934    35.883715
2935   456.867875
2936    44.361960
2937    44.783803
Name: gdp_per_capita, Length: 2938, dtype: float64
```

```
x = df[['gdp_per_capita']] # Feature: GDP per capita
y = df['Lifeexpectancy'] # Target: Life expectancy
from sklearn.model_selection import train_test_split # Import for splitting data
from sklearn.linear_model import LinearRegression # Import for linear regression
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
lr = LinearRegression()
lr.fit(x_train, y_train)
y_pred = lr.predict(x_test)
rmse_simple_lr = np.sqrt(mean_squared_error(y_test, y_pred))
mae_simple_lr = mean_absolute_error(y_test, y_pred)
print("\nSimple Linear Regression Results:")
print(f"RMSE: {rmse_simple_lr:.2f}")
print(f"MAE: {mae_simple_lr:.2f}")
```


```
plt.figure(figsize=(10, 6))
```



```
sns.scatterplot(x=x_test['gdp_per_capita'], y=y_pred, label='Predicted', color='orange')
plt.title('Life Expectancy vs GDP per Capita')
plt.xlabel('GDP per Capita (scaled)')
plt.ylabel('Life Expectancy (years)')
plt.legend()
```



Simple Linear Regression Results:  
 RMSE: 9.29  
 MAE: 7.63  
 <matplotlib.legend.Legend at 0x256c2223910>



```
# random forest regression to predict life expectancy based on GDP per capita
rf_simple = RandomForestRegressor(n_estimators=100, random_state=42)
rf_simple.fit(x_train, y_train) # Train the model
y_pred_rf_simple = rf_simple.predict(x_test) # Make predictions
rmse_rf_simple = np.sqrt(mean_squared_error(y_test, y_pred_rf_simple)) # Calculate RMSE
mae_rf_simple = mean_absolute_error(y_test, y_pred_rf_simple) # Calculate MAE
print("\nRandom Forest Regression Results:")
print(f"RMSE: {rmse_rf_simple:.2f}") # Print RMSE
print(f"MAE: {mae_rf_simple:.2f}") # Print MAE
plt.figure(figsize=(10, 6))
sns.scatterplot(x=x_test['gdp_per_capita'], y=y_pred_rf_simple, label='Predicted', color='green')
plt.title('Life Expectancy vs GDP per Capita (Random Forest)')
plt.xlabel('GDP per Capita (scaled)')
plt.ylabel('Life Expectancy (years)')
plt.legend()
plt.show() # Show the plot
```



Random Forest Regression Results:  
RMSE: 9.77  
MAE: 7.71

