

Importing necessary libraries

```

import pandas as pd # For data manipulation and analysis
import numpy as np # For numerical operations
import matplotlib.pyplot as plt # For data visualization
import seaborn as sns # For advanced data visualization
from sklearn.model_selection import train_test_split # For splitting data into training and testing sets
from sklearn.preprocessing import StandardScaler # For feature scaling
from sklearn.linear_model import LinearRegression # Linear Regression algorithm
from sklearn.ensemble import RandomForestRegressor # Random Forest algorithm
from xgboost import XGBRegressor # XGBoost algorithm
from sklearn.svm import SVR # Support Vector Regression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score # Evaluation metrics
from sklearn.impute import SimpleImputer # For handling missing values

```

Load the dataset

Assuming the dataset is downloaded and saved as 'Life Expectancy Data.csv'

```

df = pd.read_excel('led.xlsx') # Load the dataset


```

Display basic information about the dataset

```

print("Dataset Information:")
print(df.info()) # Shows columns, non-null counts, and data types
print("\nFirst 5 rows of the dataset:")
print(df.head()) # Displays first 5 rows to understand the data structure

```



```

5  infantdeaths      2938 non-null  int64
6  Alcohol          2744 non-null  float64
7  percentageexpenditure  2938 non-null  float64
8  HepatitisB       2385 non-null  float64
9  Measles          2938 non-null  int64
10 BMI              2904 non-null  float64
11 under-fivedeaths  2938 non-null  int64
12 Polio            2919 non-null  float64
13 Totalexpenditure  2712 non-null  float64
14 Diphtheria       2919 non-null  float64
15 HIV/AIDS         2938 non-null  float64
16 GDP              2490 non-null  float64
17 Population       2286 non-null  float64
18 thinness1-19years  2904 non-null  float64
19 thinness5-9years  2904 non-null  float64
20 Incomecompositionofresources  2771 non-null  float64
21 Schooling        2775 non-null  float64
dtypes: float64(16), int64(4), object(2)
memory usage: 505.1+ KB
None

```

First 5 rows of the dataset:

	Country	Year	Status	Lifeexpectancy	AdultMortality	\
0	Afghanistan	2015	Developing	65.0	263.0	
1	Afghanistan	2014	Developing	59.9	271.0	
2	Afghanistan	2013	Developing	59.9	268.0	
3	Afghanistan	2012	Developing	59.5	272.0	
4	Afghanistan	2011	Developing	59.2	275.0	

	infantdeaths	Alcohol	percentageexpenditure	HepatitisB	Measles	...	\
0	62	0.01	71.279624	65.0	1154	...	
1	64	0.01	73.523582	62.0	492	...	
2	66	0.01	73.219243	64.0	430	...	
3	69	0.01	78.184215	67.0	2787	...	
4	71	0.01	7.097109	68.0	3013	...	

	Polio	Totalexpenditure	Diphtheria	HIV/AIDS	GDP	Population	\
0	6.0	8.16	65.0	0.1	584.259210	33736494.0	
1	58.0	8.18	62.0	0.1	612.696514	327582.0	
2	62.0	8.13	64.0	0.1	631.744976	31731688.0	
3	67.0	8.52	67.0	0.1	669.959000	3696958.0	
4	68.0	7.87	68.0	0.1	63.537231	2978599.0	

	thinness1-19years	thinness5-9years	Incomecompositionofresources	\
0	17.2	17.3	0.479	
1	17.5	17.5	0.476	
2	17.7	17.7	0.470	
3	17.9	18.0	0.463	
4	18.2	18.2	0.454	

	Schooling
0	10.1
1	10.0
2	9.9
3	9.8
4	9.5

[5 rows x 22 columns]

```

df.head()

```

	Country	Year	Status	Lifexpectancy	AdultMortality	infantdeaths	Alcohol	percentageexpenditure	HepatitisB	Measles	...	Polio	Totalexpenditure	Diphtheria	HIV/
0	Afghanistan	2015	Developing	65.0	263.0	62	0.01	71.279624	65.0	1154	...	6.0	8.16	65.0	
1	Afghanistan	2014	Developing	59.9	271.0	64	0.01	73.523582	62.0	492	...	58.0	8.18	62.0	
2	Afghanistan	2013	Developing	59.9	268.0	66	0.01	73.219243	64.0	430	...	62.0	8.13	64.0	
3	Afghanistan	2012	Developing	59.5	272.0	69	0.01	78.184215	67.0	2787	...	67.0	8.52	67.0	
4	Afghanistan	2011	Developing	59.2	275.0	71	0.01	7.097109	68.0	3013	...	68.0	7.87	68.0	

5 rows × 22 columns

▼

Data Preprocessing

Check for missing values

```
print("\nMissing values in each column:")
print(df.isnull().sum()) # Counts null values in each column
```

Missing values in each column:

Country	0
Year	0
Status	0
Lifexpectancy	10
AdultMortality	10
infantdeaths	0
Alcohol	194
percentageexpenditure	0
HepatitisB	553
Measles	0
BMI	34
under-fivedeaths	0
Polio	19
Totalexpenditure	226
Diphtheria	19
HIV/AIDS	0
GDP	448
Population	652
thinness1-19years	34
thinness5-9years	34
Incomecompositionofresources	167
Schooling	163

dtype: int64

▼

Handle missing values

We'll use median imputation for numerical columns and mode for categorical

```
numerical_cols = df.select_dtypes(include=['float64', 'int64']).columns
categorical_cols = df.select_dtypes(include=['object']).columns
```

▼

Create imputers

```
num_imputer = SimpleImputer(strategy='median') # Replaces missing values with median
cat_imputer = SimpleImputer(strategy='most_frequent') # Replaces missing values with most frequent value
```

▼

Apply imputation

```
df[numerical_cols] = num_imputer.fit_transform(df[numerical_cols])
df[categorical_cols] = cat_imputer.fit_transform(df[categorical_cols])
```

▼

Verify no missing values remain

```
print("\nMissing values after imputation:")
print(df.isnull().sum())
```

Missing values after imputation:

Country	0
Year	0
Status	0
Lifexpectancy	0
AdultMortality	0
infantdeaths	0
Alcohol	0
percentageexpenditure	0
HepatitisB	0
Measles	0
BMI	0
under-fivedeaths	0
Polio	0
Totalexpenditure	0
Diphtheria	0
HIV/AIDS	0
GDP	0
Population	0
thinness1-19years	0
thinness5-9years	0
Incomecompositionofresources	0

✖ Exploratory Data Analysis (EDA)

Summary statistics

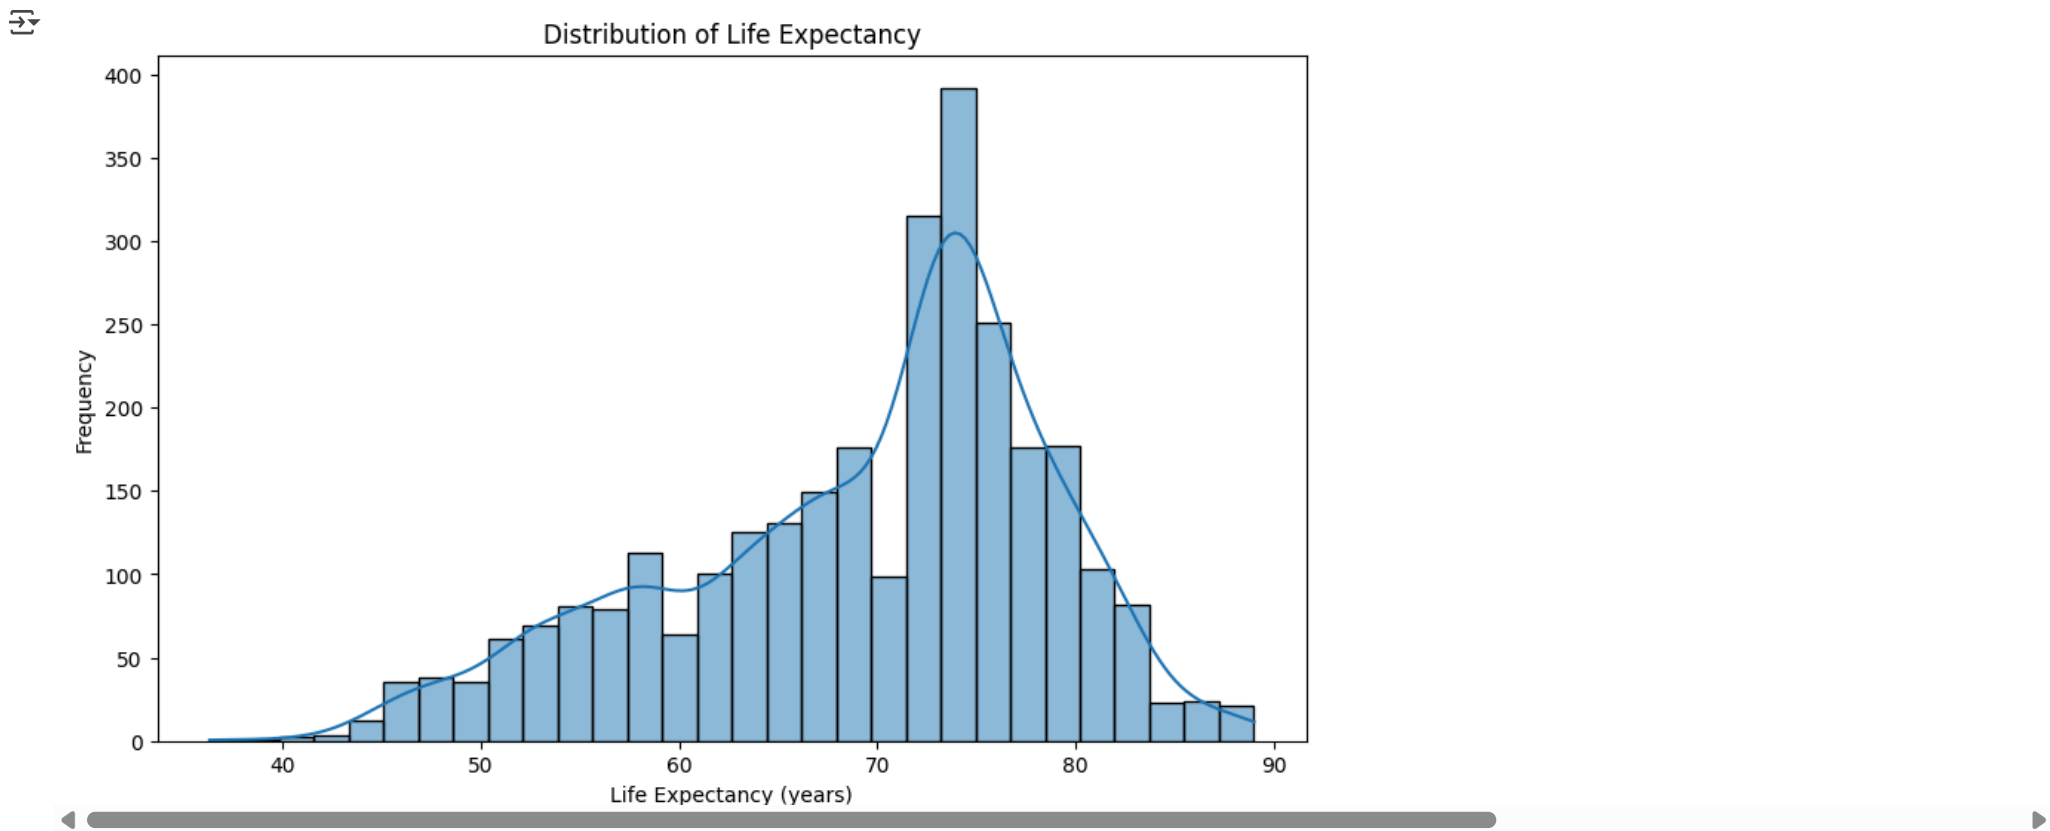
```
print("\nSummary statistics:")
print(df.describe()) # Shows count, mean, std, min, max, etc.
```

↻

Summary statistics:					
	Year	Lifeexpectancy	AdultMortality	infantdeaths	Alcohol \
count	2938.000000	2938.000000	2938.000000	2938.000000	2938.000000
mean	2007.518720	69.234717	164.725664	30.303948	4.546875
std	4.613841	9.509115	124.086215	117.926501	3.921946
min	2000.000000	36.300000	1.000000	0.000000	0.010000
25%	2004.000000	63.200000	74.000000	0.000000	1.092500
50%	2008.000000	72.100000	144.000000	3.000000	3.755000
75%	2012.000000	75.600000	227.000000	22.000000	7.390000
max	2015.000000	89.000000	723.000000	1800.000000	17.870000
	percentageexpenditure	HepatitisB	Measles	BMI \	
count	2938.000000	2938.000000	2938.000000	2938.000000	
mean	738.251295	83.022124	2419.592240	38.381178	
std	1987.914858	22.996984	11467.272489	19.935375	
min	0.000000	1.000000	0.000000	1.000000	
25%	4.685343	82.000000	0.000000	19.400000	
50%	64.912906	92.000000	17.000000	43.500000	
75%	441.534144	96.000000	360.250000	56.100000	
max	19479.911610	99.000000	212183.000000	87.300000	
	under-fivedeaths	Polio	Totalexpenditure	Diphtheria \	
count	2938.000000	2938.000000	2938.000000	2938.000000	
mean	42.035739	82.617767	5.924098	82.393125	
std	160.445548	23.367166	2.400770	23.655562	
min	0.000000	3.000000	0.370000	2.000000	
25%	0.000000	78.000000	4.370000	78.000000	
50%	4.000000	93.000000	5.755000	93.000000	
75%	28.000000	97.000000	7.330000	97.000000	
max	2500.000000	99.000000	17.600000	99.000000	
	HIV/AIDS	GDP	Population	thinness1-19years \	
count	2938.000000	2938.000000	2.938000e+03	2938.000000	
mean	1.742103	6611.523863	1.023085e+07	4.821886	
std	5.077785	13296.603449	5.402242e+07	4.397621	
min	0.100000	1.681350	3.400000e+01	0.100000	
25%	0.100000	580.486996	4.189172e+05	1.600000	
50%	0.100000	1766.947595	1.386542e+06	3.300000	
75%	0.800000	4779.405190	4.584371e+06	7.100000	
max	50.600000	119172.741800	1.293859e+09	27.700000	
	thinness5-9years	Incomecompositionofresources	Schooling		
count	2938.000000	2938.000000	2938.000000		
mean	4.852144	0.630362	12.009837		
std	4.485854	0.205140	3.265139		
min	0.100000	0.000000	0.000000		
25%	1.600000	0.504250	10.300000		
50%	3.300000	0.677000	12.300000		
75%	7.200000	0.772000	14.100000		
max	28.600000	0.948000	20.700000		

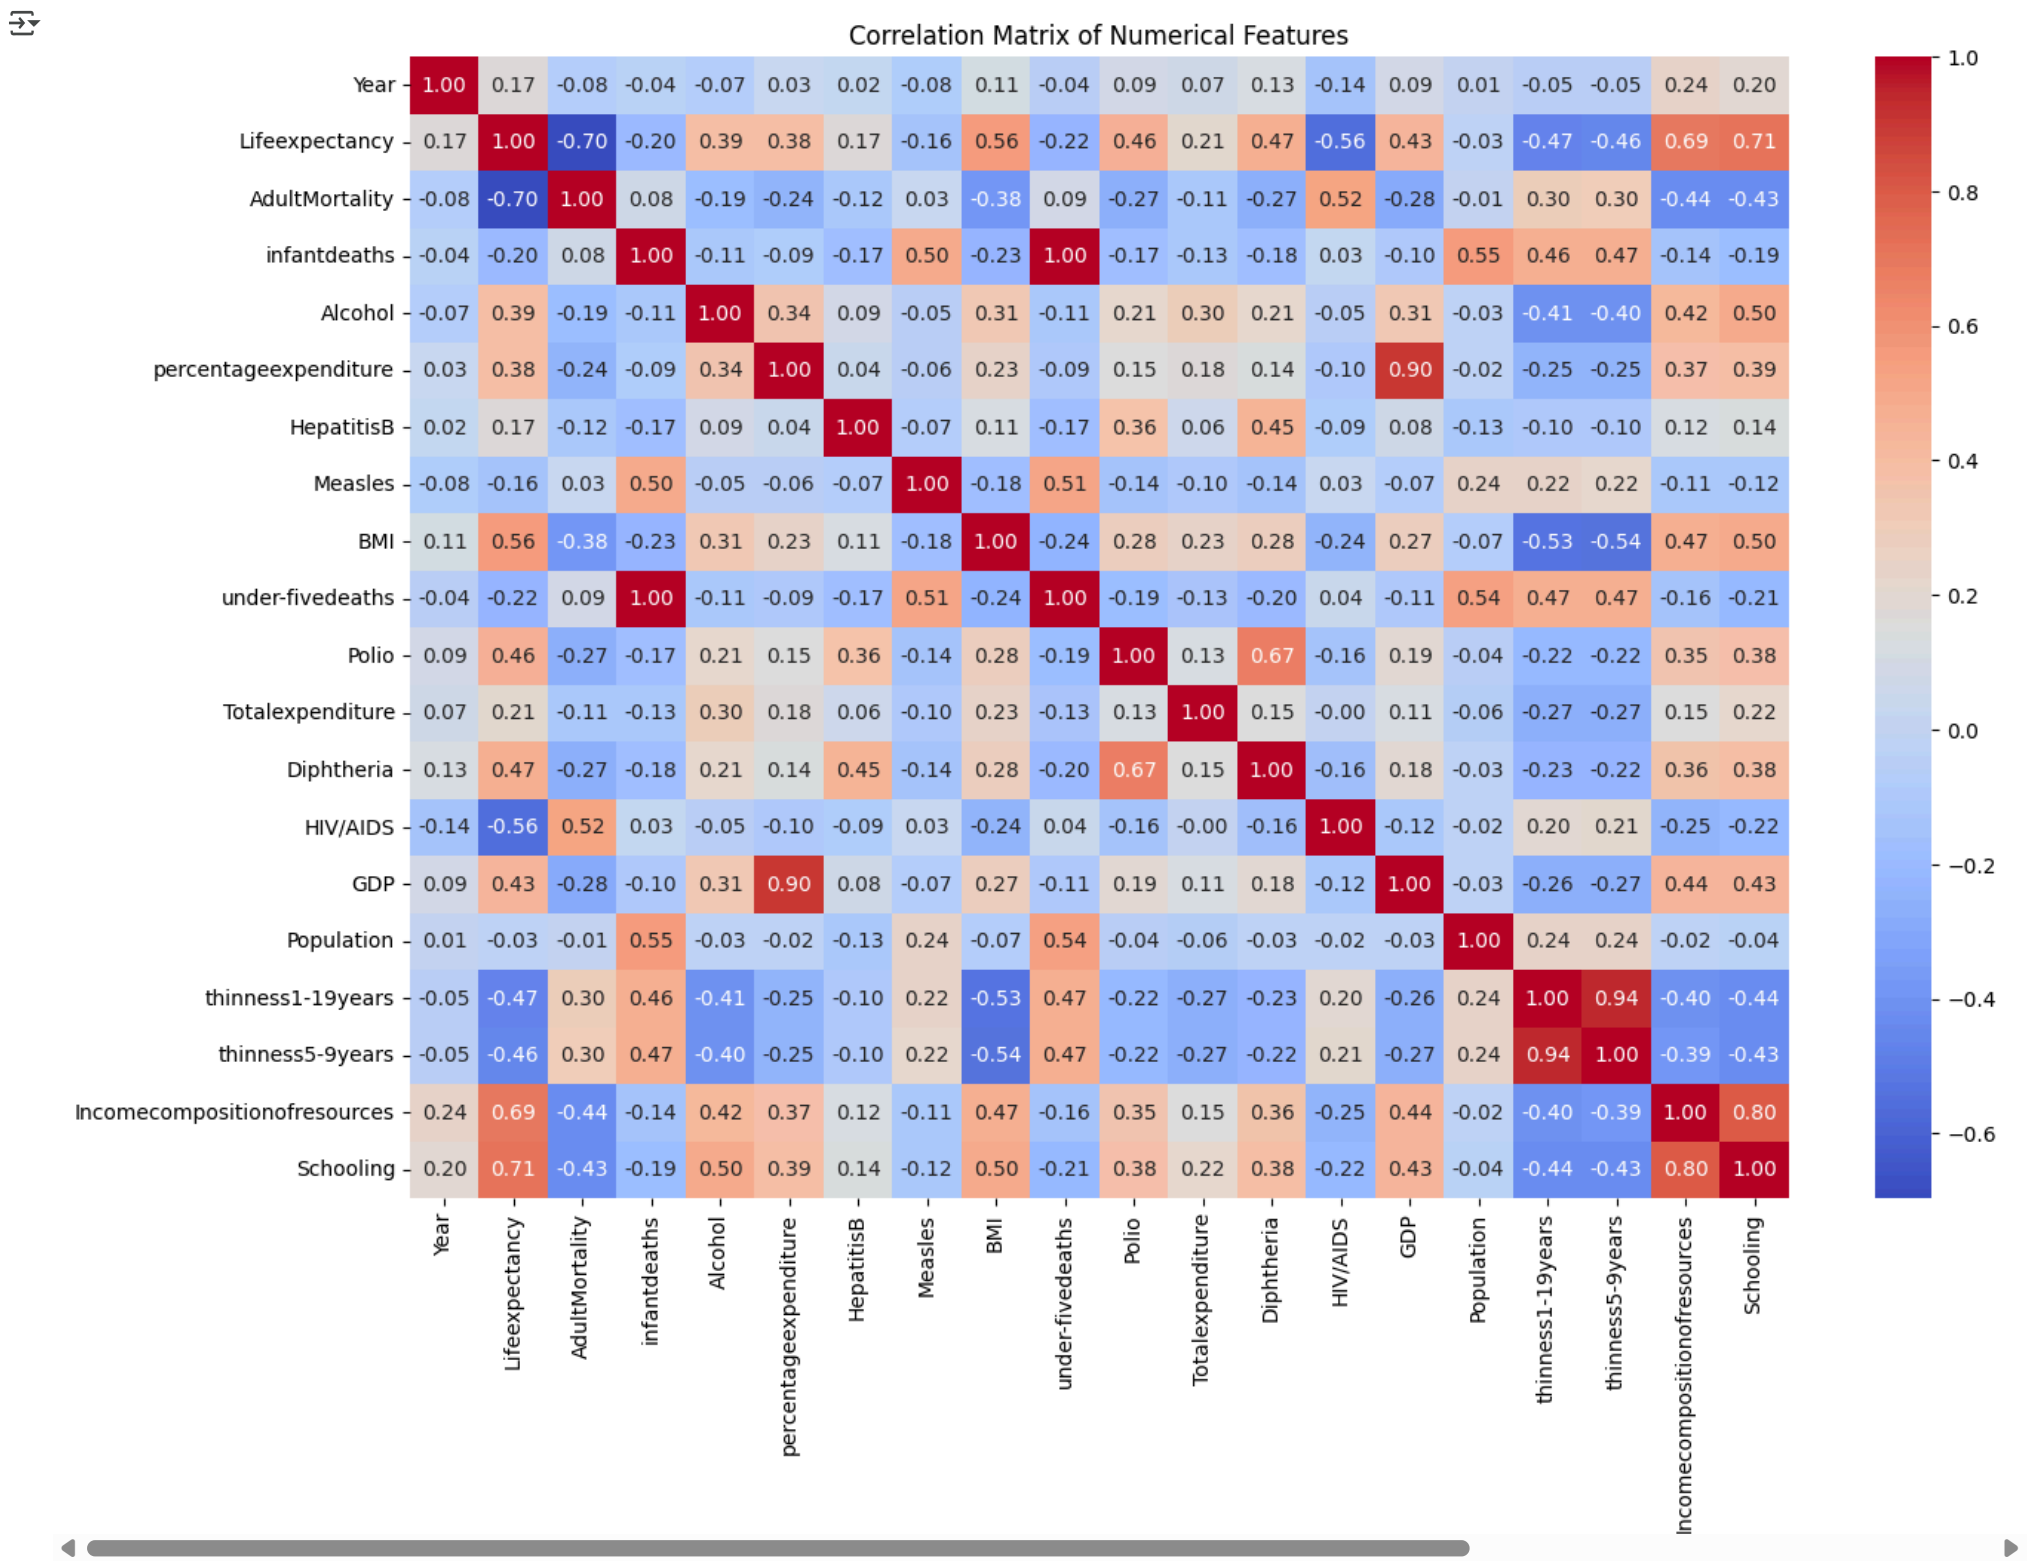
✖ Visualize the distribution of the target variable (Life expectancy)

```
plt.figure(figsize=(10, 6))
sns.histplot(df['Lifeexpectancy'],kde=True, bins=30)
plt.title('Distribution of Life Expectancy')
plt.xlabel('Life Expectancy (years)')
plt.ylabel('Frequency')
plt.show()
```










✖ Correlation matrix to understand relationships between variables








```
plt.figure(figsize=(15, 10))
corr_matrix = df.corr(numeric_only=True) # Calculates correlation between numerical columns
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix of Numerical Features')
plt.show()
```




```
#plot scatter plots for each feature against Life Expectancy
for col in numerical_cols:
    if col != 'Lifeexpectancy':
        plt.figure(figsize=(5, 3))
        sns.scatterplot(x=df[col], y=df['Lifeexpectancy'])
        plt.title(f'Scatter Plot of {col} vs Life Expectancy')
        plt.xlabel(col)
        plt.ylabel('Life Expectancy (years)')
        plt.show()
```

[4]







Scatter Plot of GDP vs Life Expectancy




Scatter Plot of Population vs Life Expectancy



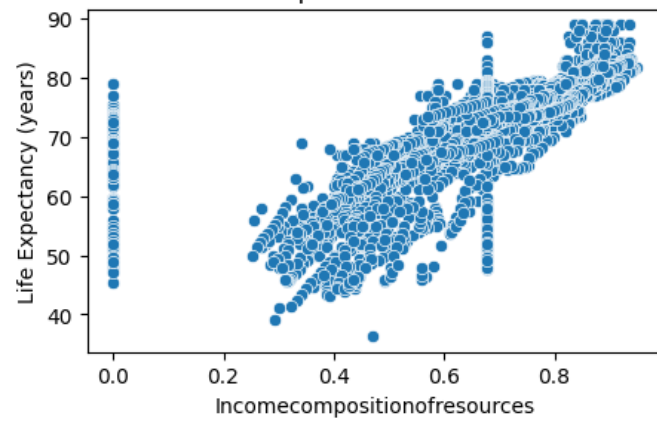
Scatter Plot of thinness1-19years vs Life Expectancy




Scatter Plot of thinness5-9years vs Life Expectancy



Scatter Plot of Incomecompositionofresources vs Life Expectancy



Scatter Plot of Schooling vs Life Expectancy



Double-click (or enter) to edit

Feature Engineering

Encoding categorical variables (Country and Status)

```
df = pd.get_dummies(df, columns=['Status'], drop_first=True) # Converts Status to binary (1 for Developed, 0 for Developing)
```

For Country, since there are many unique values, we'll drop it to avoid high dimensionality

```
df.drop('Country', axis=1, inplace=True) # Removes Country column
```

Define features (X) and target (y)

```
X = df.drop('Lifeexpectancy', axis=1) # All columns except Life expectancy
y = df['Lifeexpectancy'] # Target variable we want to predict
```

Split the data into training and testing sets (80% train, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Feature Scaling

Many algorithms perform better when features are on similar scales

```
scaler = StandardScaler() # Standardizes features by removing mean and scaling to unit variance
X_train_scaled = scaler.fit_transform(X_train) # Fit to training data and transform
X_test_scaled = scaler.transform(X_test) # Transform test data using same scaling
```

Model Building and Evaluation

Dictionary to store evaluation results

1. Linear Regression

```
results = {}

# 1. Linear Regression
lr = LinearRegression() # Creates a linear regression model
lr.fit(X_train_scaled, y_train) # Trains the model on scaled training data
y_pred_lr = lr.predict(X_test_scaled) # Makes predictions on test data
```

Calculate evaluation metrics

```
rmse_lr = np.sqrt(mean_squared_error(y_test, y_pred_lr)) # Root Mean Squared Error
mae_lr = mean_absolute_error(y_test, y_pred_lr) # Mean Absolute Error
r2_lr = r2_score(y_test, y_pred_lr) # R-squared score

# Store results
results['Linear Regression'] = {'RMSE': rmse_lr, 'MAE': mae_lr, 'R2': r2_lr}
```

2. Random Forest Regressor

```
rf = RandomForestRegressor(n_estimators=100, random_state=42) # Creates RF model with 100 trees
rf.fit(X_train_scaled, y_train) # Trains the model
y_pred_rf = rf.predict(X_test_scaled) # Makes predictions
```

Double-click (or enter) to edit

Calculate metrics

```
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))
mae_rf = mean_absolute_error(y_test, y_pred_rf)
r2_rf = r2_score(y_test, y_pred_rf)
```

Store results


```
results['Random Forest'] = {'RMSE': rmse_rf, 'MAE': mae_rf, 'R2': r2_rf}
```

3. XGBoost Regressor

```
xgb = XGBRegressor(n_estimators=100, random_state=42) # Creates XGBoost model with 100 trees
xgb.fit(X_train_scaled, y_train) # Trains the model
y_pred_xgb = xgb.predict(X_test_scaled) # Makes predictions
```

Calculate metrics

```
rmse_xgb = np.sqrt(mean_squared_error(y_test, y_pred_xgb))
mae_xgb = mean_absolute_error(y_test, y_pred_xgb)
r2_xgb = r2_score(y_test, y_pred_xgb)
```

Store results

```
results['XGBoost'] = {'RMSE': rmse_xgb, 'MAE': mae_xgb, 'R2': r2_xgb}
```

4. Support Vector Regression

```
svr = SVR(kernel='rbf') # Creates SVR model with Radial Basis Function kernel
svr.fit(X_train_scaled, y_train) # Trains the model
y_pred_svr = svr.predict(X_test_scaled) # Makes predictions
```

Calculate metrics

```
rmse_svr = np.sqrt(mean_squared_error(y_test, y_pred_svr))
mae_svr = mean_absolute_error(y_test, y_pred_svr)
r2_svr = r2_score(y_test, y_pred_svr)
```

Store results

```
results['Support Vector Regression'] = {'RMSE': rmse_svr, 'MAE': mae_svr, 'R2': r2_svr}
```

Display results

```
print("\nModel Evaluation Results:")
for model, metrics in results.items():
    print(f"\n{model}:")
    print(f"RMSE: {metrics['RMSE']:.2f}") # Lower is better
    print(f"MAE: {metrics['MAE']:.2f}") # Lower is better
    print(f"R2 Score: {metrics['R2']:.2f}") # Closer to 1 is better
```

```
↔
Model Evaluation Results:

Linear Regression:
RMSE: 3.91
MAE: 2.86
R2 Score: 0.82

Random Forest:
RMSE: 1.66
MAE: 1.08
R2 Score: 0.97

XGBoost:
RMSE: 1.75
MAE: 1.18
R2 Score: 0.96

Support Vector Regression:
RMSE: 3.32
MAE: 2.31
R2 Score: 0.87
```

Feature Importance for the best performing model (Random Forest or XGBoost)

Let's check feature importance from Random Forest

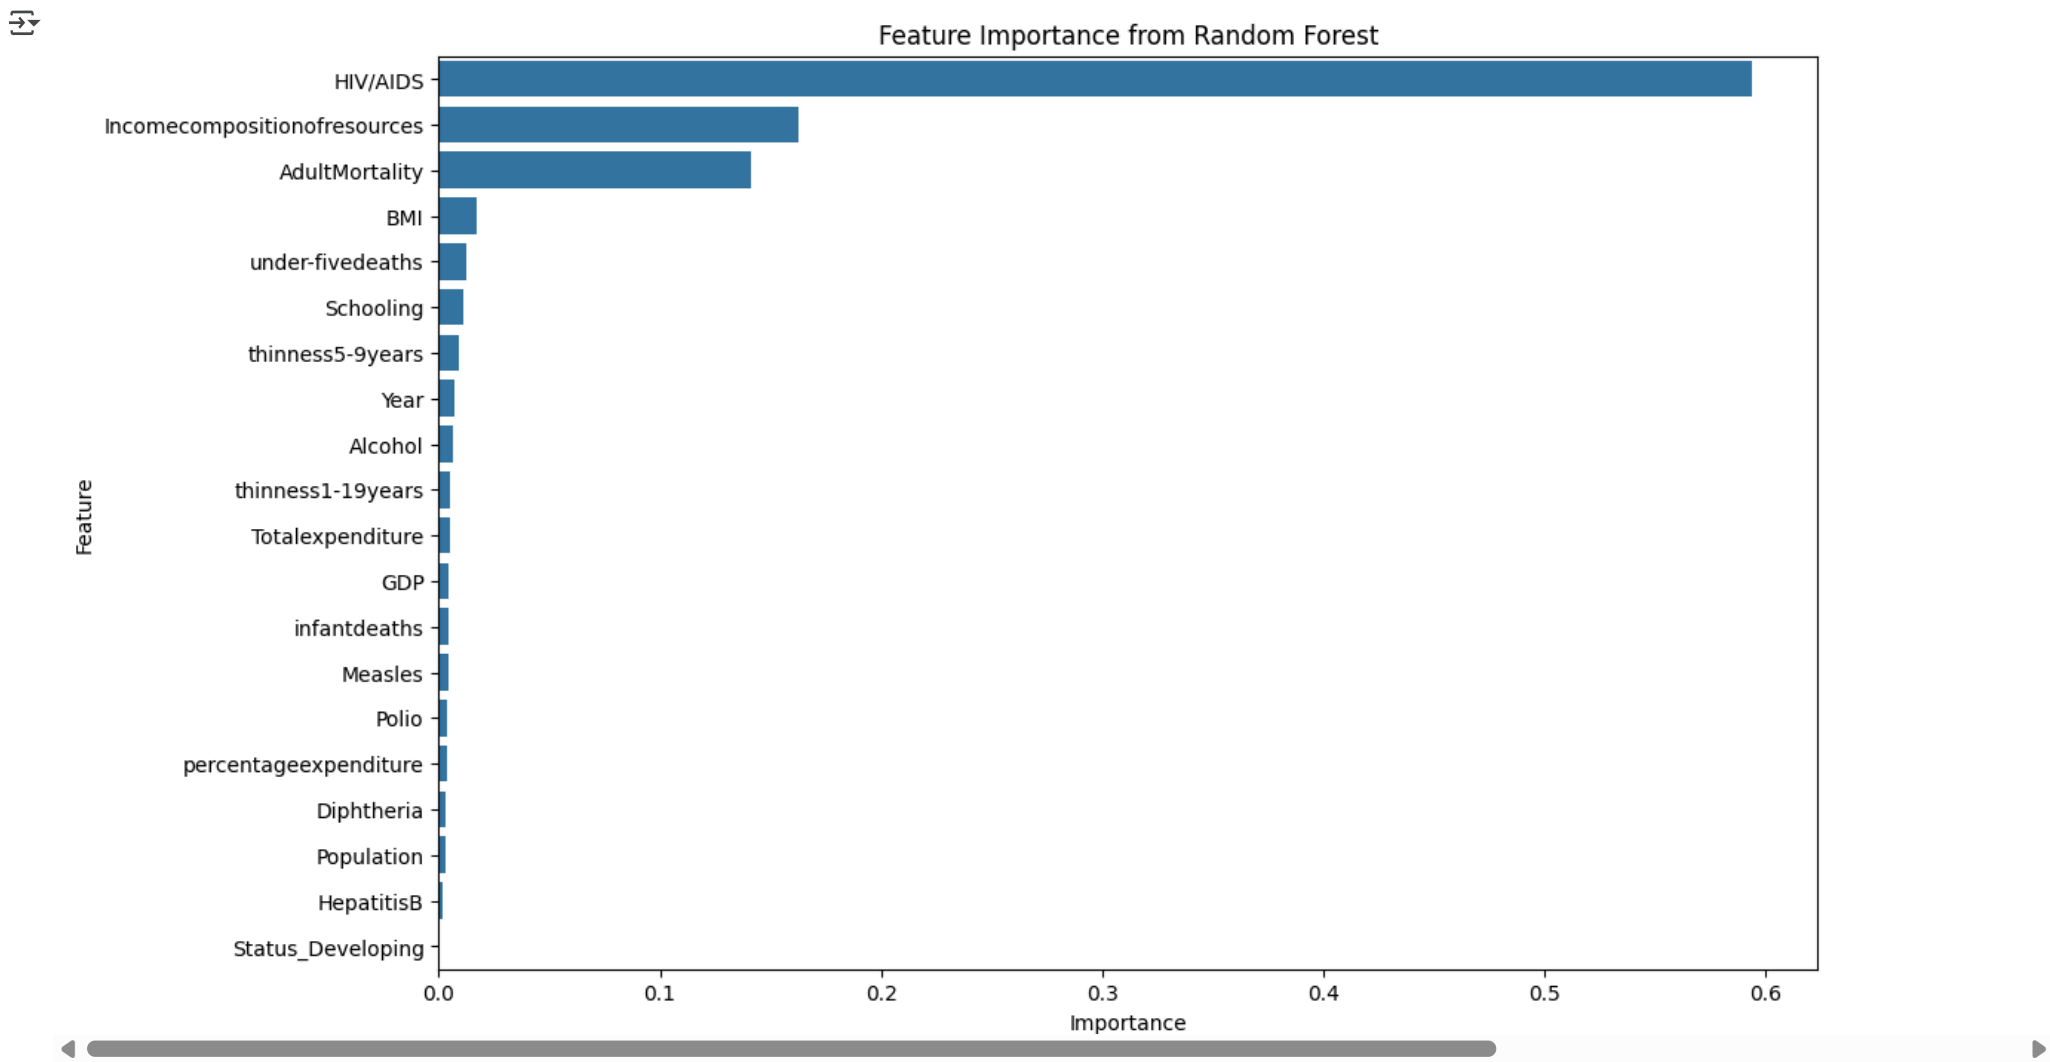
```
feature_importance = rf.feature_importances_ # Gets importance scores from trained RF model
features = X.columns # Gets feature names
```

Create a dataframe for visualization

```
importance_df = pd.DataFrame({'Feature': features, 'Importance': feature_importance})
importance_df = importance_df.sort_values('Importance', ascending=False) # Sorts by importance
```

Plot feature importance

```
plt.figure(figsize=(12, 8))
sns.barplot(x='Importance', y='Feature', data=importance_df)
plt.title('Feature Importance from Random Forest')
plt.show()
```



gdp per capita

```
#implement simple linear regression to predict life expectancy based solely on GDP per capita
df['gdp_per_capita'] = df['GDP'] / df['Population'] # Calculate GDP per capita
#MULTIPLY GDP per capita by 100000 to scale it
df['gdp_per_capita'] *= 100000 # Scale GDP per capita for better interpretability
print("\nGDP per Capita:")
print(df['gdp_per_capita']) # Display first few rows of GDP per capita
```

```
GDP per Capita:
0      17.318314
1     1870.360746
2      19.908962
3     181.218991
4      21.331247
...
2933    35.559872
2934    35.883715
2935    456.867875
2936     44.361960
2937     44.783803
Name: gdp_per_capita, Length: 2938, dtype: float64
```


```
x = df[['gdp_per_capita']] # Feature: GDP per capita
y = df['Lifeexpectancy'] # Target: Life expectancy
from sklearn.model_selection import train_test_split # Import for splitting data
from sklearn.linear_model import LinearRegression # Import for linear regression
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
lr = LinearRegression()
lr.fit(x_train, y_train)
y_pred = lr.predict(x_test)
rmse_simple_lr = np.sqrt(mean_squared_error(y_test, y_pred))
mae_simple_lr = mean_absolute_error(y_test, y_pred)
print("\nSimple Linear Regression Results:")
print(f"RMSE: {rmse_simple_lr:.2f}")
print(f"MAE: {mae_simple_lr:.2f}")
```

```
plt.figure(figsize=(10, 6))

#sns.scatterplot(x=x_test['gdp_per_capita'], y=y_pred, label='Predicted', color='orange')
sns.scatterplot(x=x_test['gdp_per_capita'], y=y_test, label='Actual', color='blue')
plt.title('Life Expectancy vs GDP per Capita')
plt.xlabel('GDP per Capita (scaled)')
plt.ylabel('Life Expectancy (years)')
plt.legend()
```




Simple Linear Regression Results:
RMSE: 9.29
MAE: 7.63
<matplotlib.legend.Legend at 0x23aa774b1d0>




```
# random forest regression to predict life expectancy based on GDP per capita
rf_simple = RandomForestRegressor(n_estimators=100, random_state=42)
rf_simple.fit(x_train, y_train) # Train the model
y_pred_rf_simple = rf_simple.predict(x_test) # Make predictions
rmse_rf_simple = np.sqrt(mean_squared_error(y_test, y_pred_rf_simple)) # Calculate RMSE
mae_rf_simple = mean_absolute_error(y_test, y_pred_rf_simple) # Calculate MAE
print("\nRandom Forest Regression Results:")
print(f"RMSE: {rmse_rf_simple:.2f}") # Print RMSE
print(f"MAE: {mae_rf_simple:.2f}") # Print MAE
plt.figure(figsize=(10, 6))
sns.scatterplot(x=x_test['gdp_per_capita'], y=y_pred_rf_simple, label='Predicted', color='green')
plt.title('Life Expectancy vs GDP per Capita (Random Forest)')
plt.xlabel('GDP per Capita (scaled)')
plt.ylabel('Life Expectancy (years)')
plt.legend()
plt.show() # Show the plot
```



Random Forest Regression Results:
RMSE: 9.77
MAE: 7.71



```
sns.scatterplot(x=y_test, y=y_pred, label='Predicted (Linear Regression)', color='orange')
```



df

	Year	Lifeexpectancy	AdultMortality	infantdeaths	Alcohol	percentageexpenditure	HepatitisB	Measles	BMI	under-fivedeaths	...	Diphtheria	HIV/AIDS	GDP	Populati
0	2015.0	65.0	263.0	62.0	0.01	71.279624	65.0	1154.0	19.1	83.0	...	65.0	0.1	584.259210	33736494
1	2014.0	59.9	271.0	64.0	0.01	73.523582	62.0	492.0	18.6	86.0	...	62.0	0.1	612.696514	327584
2	2013.0	59.9	268.0	66.0	0.01	73.219243	64.0	430.0	18.1	89.0	...	64.0	0.1	631.744976	31731688
3	2012.0	59.5	272.0	69.0	0.01	78.184215	67.0	2787.0	17.6	93.0	...	67.0	0.1	669.959000	3696958
4	2011.0	59.2	275.0	71.0	0.01	7.097109	68.0	3013.0	17.2	97.0	...	68.0	0.1	63.537231	2978598
...
2933	2004.0	44.3	723.0	27.0	4.36	0.000000	68.0	31.0	27.1	42.0	...	65.0	33.6	454.366654	12777511
2934	2003.0	44.5	715.0	26.0	4.06	0.000000	7.0	998.0	26.7	41.0	...	68.0	36.7	453.351155	12633897
2935	2002.0	44.8	73.0	25.0	4.43	0.000000	73.0	304.0	26.3	40.0	...	71.0	39.8	57.348340	125521
2936	2001.0	45.3	686.0	25.0	1.72	0.000000	76.0	529.0	25.9	39.0	...	75.0	42.1	548.587312	12366168
2937	2000.0	46.0	665.0	24.0	1.68	0.000000	79.0	1483.0	25.5	39.0	...	78.0	43.5	547.358878	12222251

2938 rows × 22 columns

Create a linear regression model to predict life expectancy from immunization factors (Polio, Diphtheria, Hepatitis B). Interpret the coefficients - which vaccine has the strongest association with life expectancy? Use Data set of life expectancy shared earlier .

#Create a linear regression model to predict life expectancy from immunization factors (Polio, Diphtheria, Hepatitis B). Interpret the coefficients - which vaccine has the stronge

```
immunization_features1 = df[['Schooling', 'Incomecompositionofresources']]
immunization_features = df[['Polio', 'Diphtheria', 'HepatitisB']]
immunization_target = df['Lifeexpectancy']
immunization_X_train, immunization_X_test, immunization_y_train, immunization_y_test = train_test_split(
    immunization_features1, immunization_target, test_size=0.2, random_state=42
)
```

```
from sklearn.linear_model import LinearRegression
immunization_lr = LinearRegression()
immunization_lr.fit(immunization_X_train, immunization_y_train)
immunization_y_pred = immunization_lr.predict(immunization_X_test)
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
immunization_rmse = np.sqrt(mean_squared_error(immunization_y_test, immunization_y_pred))
immunization_mae = mean_absolute_error(immunization_y_test, immunization_y_pred)
immunization_r2 = r2_score(immunization_y_test, immunization_y_pred)
print("\nImmunization Linear Regression Results:")
print(f"RMSE: {immunization_rmse:.2f}")
print(f"MAE: {immunization_mae:.2f}")
print(f"R2 Score: {immunization_r2:.2f}")
```

Immunization Linear Regression Results:
RMSE: 6.13
MAE: 4.33
R2 Score: 0.57

Start coding or [generate](#) with AI.

```
coefficients = immunization_lr.coef_
intercept = immunization_lr.intercept_

print("\nImmunization Linear Regression Coefficients:")
immunization_feature_names = immunization_features1.columns
print("\nImmunization Feature Coefficients:")
for feature, coef in zip(immunization_feature_names, coefficients):
    print(f"{feature}: {coef:.4f}")

print("\nImmunization Feature INTERCEPTS:",intercept)

# Interpretation of coefficients
print("\nInterpretation of Coefficients:")
for feature, coef in zip(immunization_feature_names, coefficients):
    if coef > 0:
        print(f"A unit increase in {feature} is associated with an increase in life expectancy by {coef:.4f} years.")
```

```
else:
    print(f"A unit increase in {feature} is associated with a decrease in life expectancy by {-coef:.4f} years.")
```


Immunization Linear Regression Coefficients:

Immunization Feature Coefficients:

Schooling: 1.2522

Incomecompositionofresources: 16.5652

Immunization Feature INTERCEPTS: 43.710149059244344

Interpretation of Coefficients:

A unit increase in Schooling is associated with an increase in life expectancy by 1.2522 years.

A unit increase in Incomecompositionofresources is associated with an increase in life expectancy by 16.5652 years.

```
plt.figure(figsize=(10, 6))
```

```
sns.scatterplot(x=immunization_y_test, y=immunization_y_pred, label='Predicted', color='orange')
```

```
plt.plot(immunization_y_test, immunization_y_test, color='red', linestyle='--', label='Actual') # Diagonal line for actual values
```

```
plt.grid(True)
```

```
plt.grid(True)
plt.title('Life Expectancy vs Immunization Factors')
```


```
plt.xlabel('Actual Life Expectancy (years)')
```

```
plt.xlabel('Actual Life Expectancy (years)')
plt.ylabel('Predicted Life Expectancy (years)')
```

```
plt.legend()
```



```
<matplotlib.legend.Legend at 0x23ab2ad0a50>
```



• •

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.